VISVESVARAYA TECHNOLOGICAL UNIVERSITY

"JnanaSangama", Belgaum- 590014, Karnataka.

LAB REPORT on

Machine Learning (23CS6PCMAL)

Submitted by

Pranav R Hegde (1BM22CS202)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING

B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU - 560019

February 2025 – June 25

B.M.S. College of Engineering

Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)

Department of Computer Science and Engineering

CERTIFICATE

This is to certify that the Lab work entitled "Machine Learning (23CS6PCMAL)" carried out by **Pranav Hegde (1BM22CS202),** who is bonafide student of **B.M.S. College of Engineering.** It is in partial fulfilment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Laboratory report has been approved as it satisfies the academic requirements in respect of a Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Lab Faculty Incharge Name: **Ms. Saritha A N** Assistant Professor Department of CSE, BMSCE	 **Dr. Kavitha Sooda** Professor & HOD Department of CSE, BMSCE

INDEX

11	12-05-2025	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	44

Github Link: https://github.com/pranav0hegde/ML

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot:

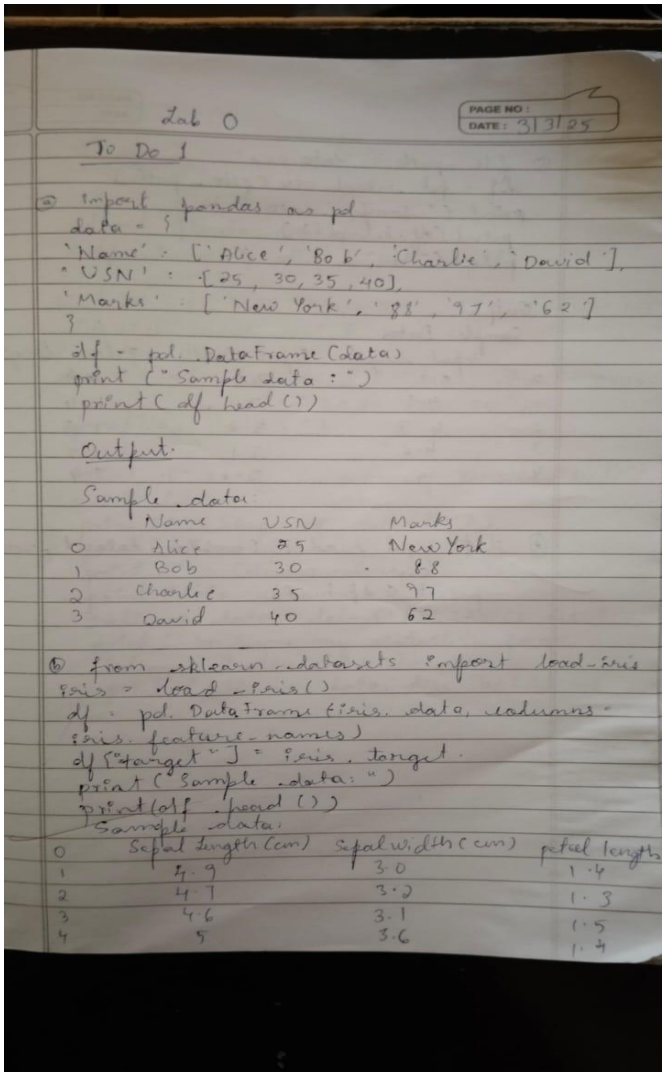

Lab 0

PAGE NO:
DATE: 3/3/25

To Do 1

① import pandas as pd
data = {
'Name': ['Alice', 'Bob', 'Charlie', 'David'],
'USN': [25, 30, 35, 40],
'Marks': ['New York', '88', '97', '62']
}
df = pd.DataFrame(data)
print("Sample data:")
print(df.head())

Output:

Sample data:

	Name	USN	Marks
0	Alice	25	New York
1	Bob	30	88
2	Charlie	35	97
3	David	40	62

② from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df["target"] = iris.target
print("Sample data:")
print(df.head())
Sample data:

	Sepal length (cm)	Sepal width (cm)	petal length
0			
1	4.9	3.0	1.4
2	4.7	3.2	1.3
3	4.6	3.1	1.5
4	5	3.6	1.4

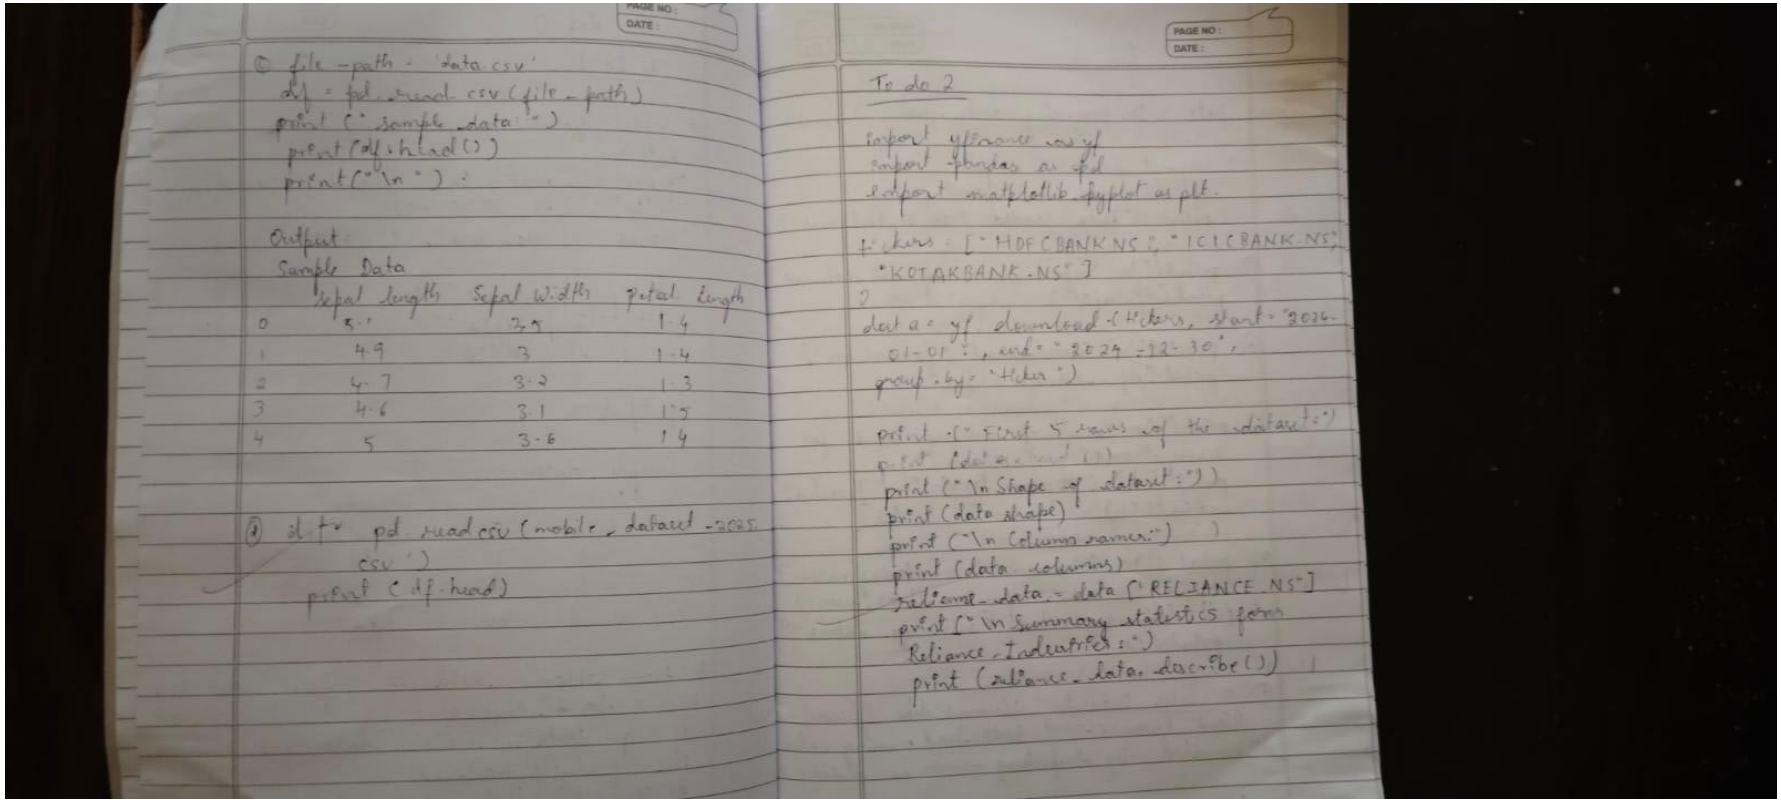

```
file_path = 'data.csv'
df = pd.read_csv(file_path)
print("Sample data:")
print(df.head())
print("\n")
```

Output:

Sample Data

	Sepal length	Sepal Width	Petal Length
0	5.1	3.5	1.4
1	4.9	3	1.4
2	4.7	3.2	1.3
3	4.6	3.1	1.5
4	5	3.6	1.4

```
df = pd.read_csv('mobile_dataset_2025.csv')
print(df.head)
```

To do 2

```
import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt

tickers = ["HDFCBANK.NS", "ICICIBANK.NS",
"KOTAKBANK.NS"]
?
data = yf.download(tickers, start="2024-
01-01", end="2024-12-30",
group_by='ticker')

print("First 5 rows of the dataset:")
print(data.head())
print("\n Shape of dataset:")
print(data.shape)
print("\n Column names:")
print(data.columns)
reliance_data = data["RELIANCE.NS"]
print("\n Summary statistics for
Reliance Industries:")
print(reliance_data.describe())
```

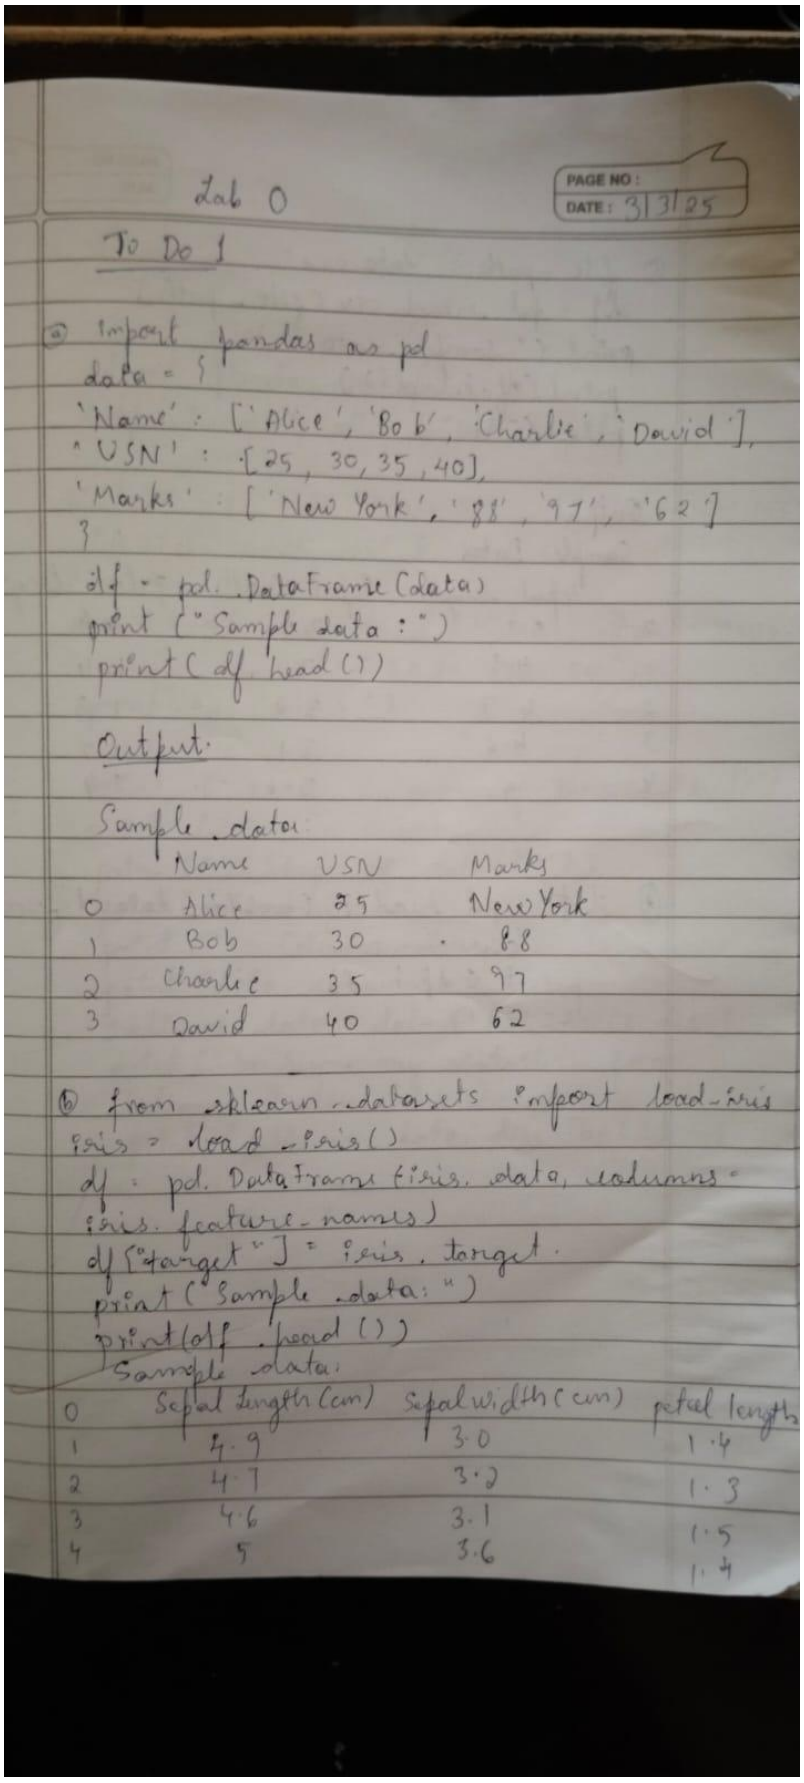

Lab 0

PAGE NO :
DATE : 3/3/25

To Do 1

① import pandas as pd
data = {
'Name' : ['Alice', 'Bob', 'Charlie', 'David'],
'USN' : [25, 30, 35, 40],
'Marks' : ['New York', '88', '97', '62']
}
df = pd.DataFrame(data)
print("Sample data :")
print(df.head())

Output:

Sample data:

	Name	USN	Marks
0	Alice	25	New York
1	Bob	30	88
2	Charlie	35	97
3	David	40	62

② from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=
iris.feature_names)
df["target"] = iris.target
print("Sample data:")
print(df.head())
Sample data:

	Sepal Length (cm)	Sepal width (cm)	petal length
0			
1	4.9	3.0	1.4
2	4.7	3.2	1.3
3	4.6	3.1	1.5
4	5	3.6	1.4

Code:

```
from sklearn.datasets import load_iris import

pandas as pd

iris = load_iris()

df = pd.DataFrame(iris.data, columns=iris.feature_names)

df.head()


df['target'] = iris.target

df

import kagglehub



# Download latest version

path = kagglehub.dataset_download("abdulmalik1518/mobiles-dataset-2025")



print("Path to dataset files:", path)



df = pd.read_csv("/content/Mobiles_Dataset_(2025).csv", encoding='latin-1') # or 'ISO-8859-1', or 'cp1252'

df.head()

df['Company Name']



data = {"USN" : ['1', "2", "3"], "Name" : ["A", "B", "C"]}

df = pd.DataFrame(data)

df
```

```
from sklearn.datasets import load_diabetes diabetes =
load_diabetes() df = pd.DataFrame(diabetes.data,
columns=diabetes.feature_names) df.head()
df.columns
df = pd.read_csv("/content/Dataset_of_Diabetes .csv")
df.head() import yfinance as yf import pandas as pd
import matplotlib.pyplot as plt



tickers = ["RELIANCE.NS", "TCS.NS", "INFY.NS"]
# Fetch historical data for the last 1 year
data = yf.download(tickers, start="2022-10-01", end="2023-10-01", group_by='ticker')


# Display the first 5 rows of the dataset
print("First 5 rows of the dataset:")


print(data.head())
print("\nShape of the dataset:")
print(data.shape)
# Summary statistics for a specific stock (e.g., Reliance)
reliance_data = data['RELIANCE.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
```

```
# Calculate daily returns

reliance_data['Daily Return'] = reliance_data['Close'].pct_change()


# Plot the closing price and daily returns

plt.figure(figsize=(12, 6))

plt.subplot(2, 1, 1)

reliance_data['Close'].plot(title="Reliance Industries - Closing Price")

plt.subplot(2, 1, 2)

reliance_data['Daily Return'].plot(title="Reliance Industries - Daily Returns",
color='orange') plt.tight_layout() plt.show()
```

Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:

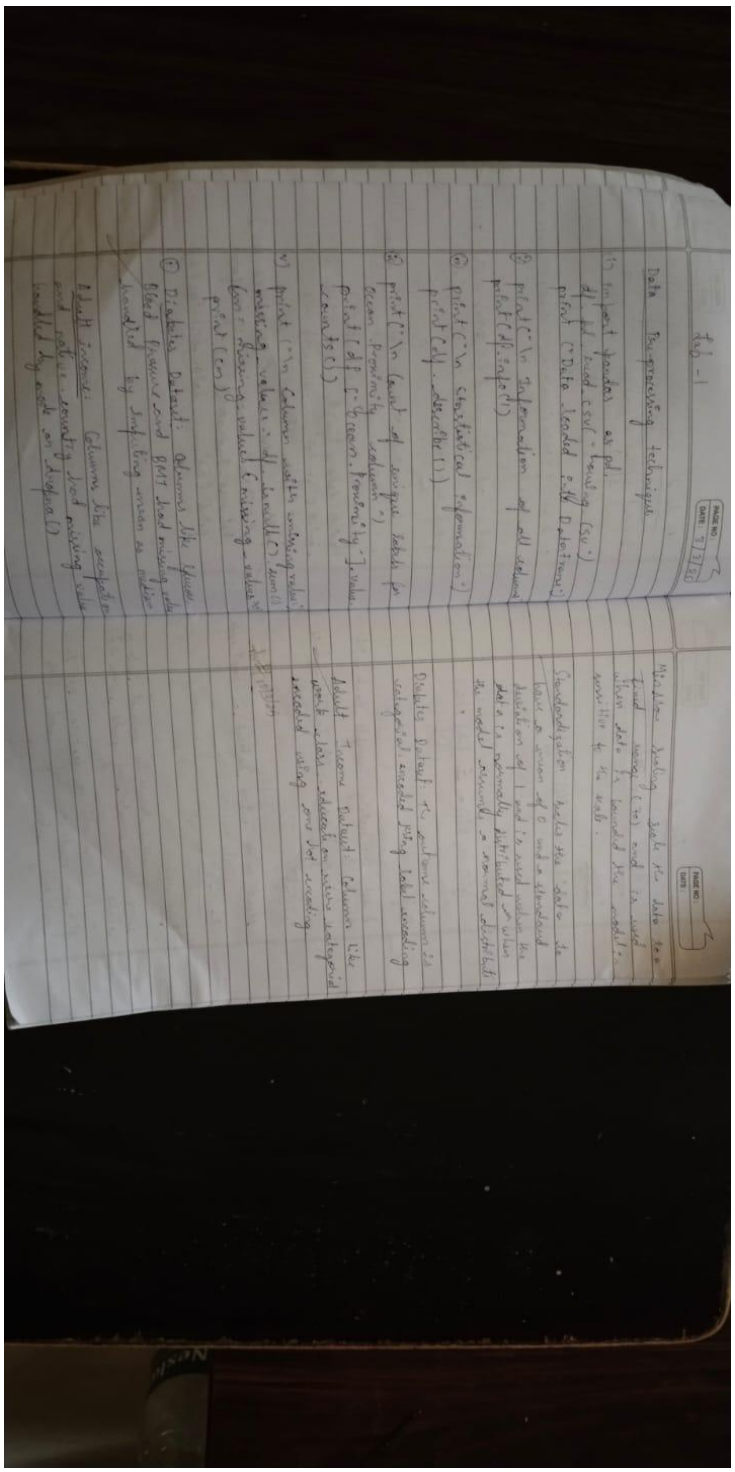

Code:

```
import pandas as pd import
numpy as np

# Load dataset
df = pd.read_csv("data.csv")
print(df.head())
# Check missing values
print(df.isnull().sum())

# Drop rows with missing values df_cleaned
= df.dropna()

# Or fill missing values with mean/median
df['Age'].fillna(df['Age'].mean(), inplace=True)
df['Salary'].fillna(df['Salary'].median(),
inplace=True)

# For nominal categories
df = pd.get_dummies(df, columns=['Gender', 'Country'], drop_first=True)

# For ordinal categories
```

```
from sklearn.preprocessing import OrdinalEncoder encoder
= OrdinalEncoder()
df[['Education_Level']] = encoder.fit_transform(df[['Education_Level']])


from sklearn.preprocessing import StandardScaler, MinMaxScaler

# Standardization (Z-score) scaler
= StandardScaler() df[['Age',
'Salary']] =
scaler.fit_transform(df[['Age',
'Salary']])

# Min-Max Normalization minmax
= MinMaxScaler()
df[['Age', 'Salary']] = minmax.fit_transform(df[['Age', 'Salary']])


# Using IQR method
Q1 = df['Salary'].quantile(0.25)
Q3 = df['Salary'].quantile(0.75)
IQR = Q3 - Q1
```

```
df = df[(df['Salary'] >= Q1 - 1.5*IQR) & (df['Salary'] <= Q3 + 1.5*IQR)]



df['Age_Salary_Ratio'] = df['Age'] / df['Salary']



# Drop irrelevant columns

df.drop(['User_ID', 'Name'], axis=1, inplace=True)


# Correlation-based filtering correlation_matrix

= df.corr() print(correlation_matrix)


from sklearn.model_selection import train_test_split


X = df.drop('Purchased', axis=1) y

= df['Purchased']


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample

Screenshot:

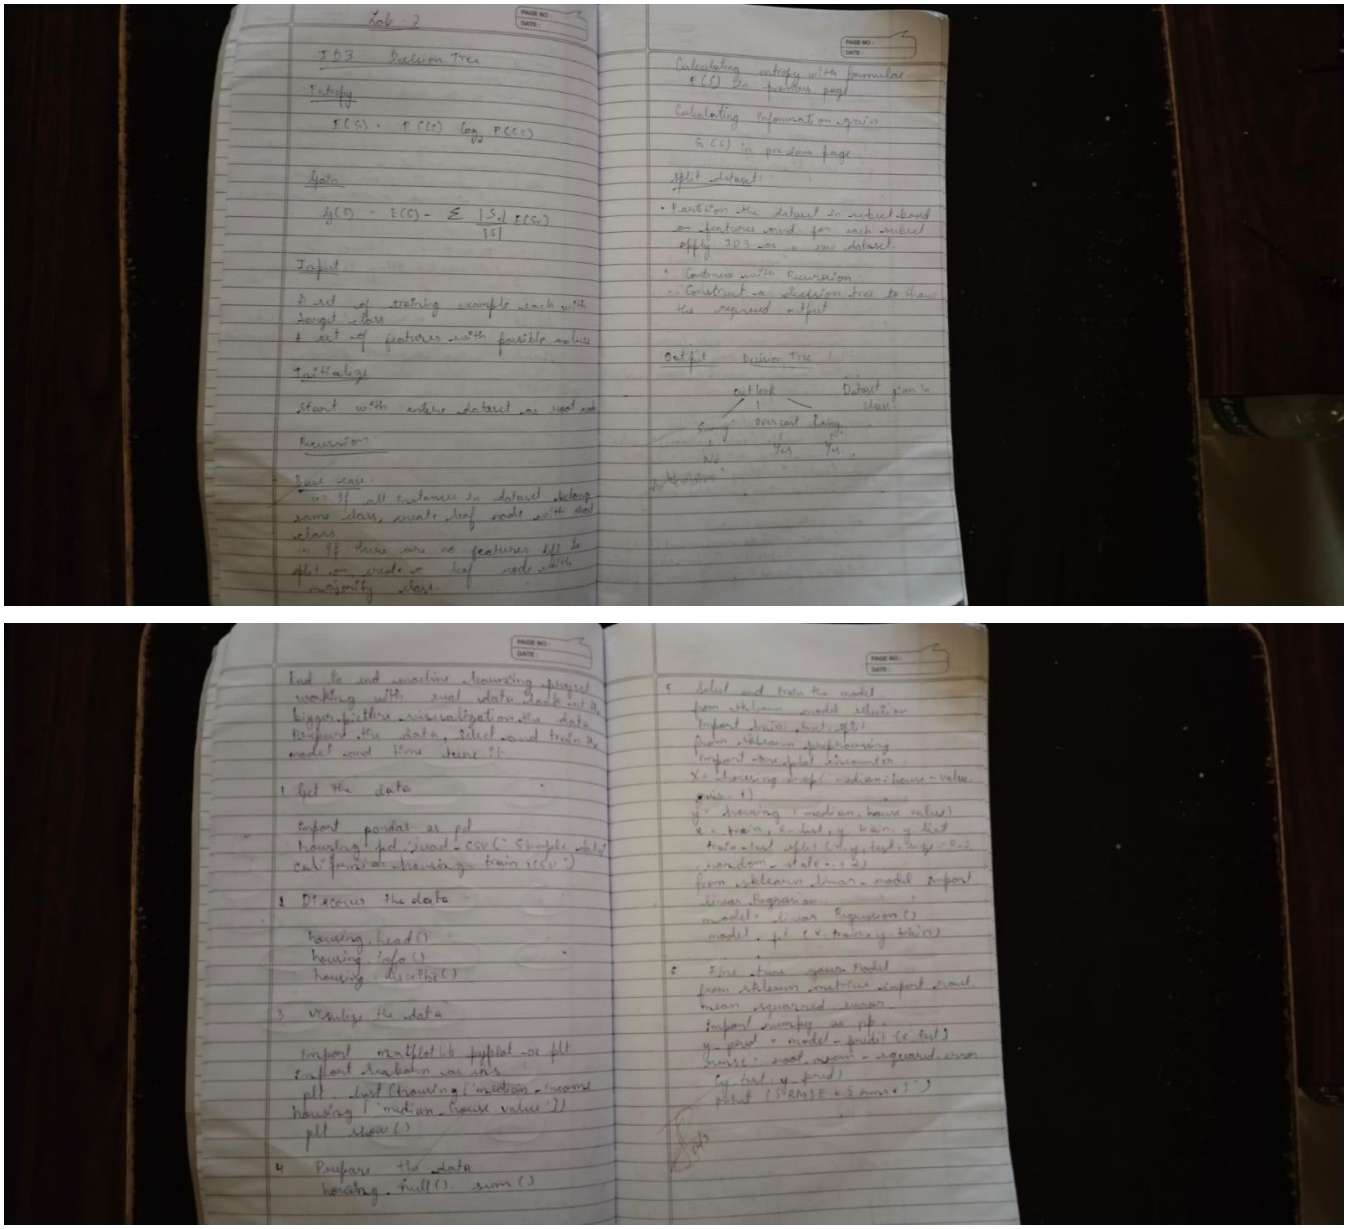

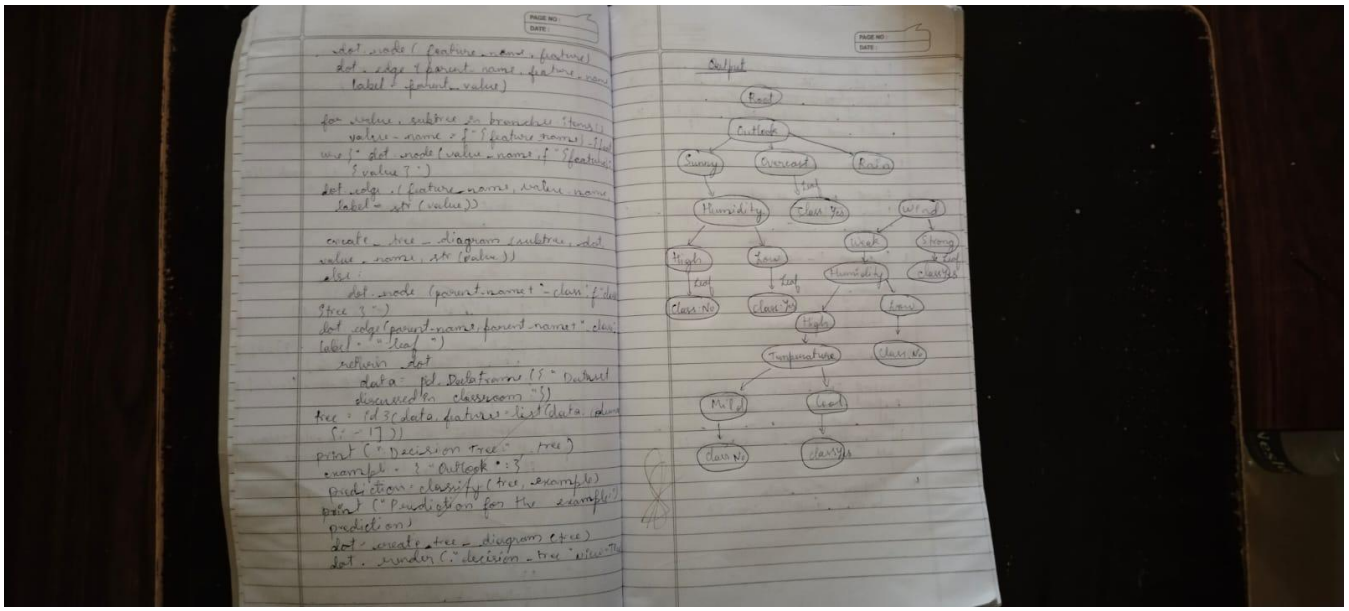

Code:

```
import pandas as pd import
numpy as np
from graphviz import Digraph


# Calculate Entropy def
entropy(data):
    class_probabilities = data.iloc[:, -1].value_counts(normalize=True) return
    -np.sum(class_probabilities * np.log2(class_probabilities))


# Calculate Information Gain
def information_gain(data, feature):
   total_entropy = entropy(data)
```

```
    feature_values = data[feature].unique()
    weighted_entropy = 0
    for value in feature_values:
        subset = data[data[feature] == value]
        weighted_entropy += (len(subset) / len(data)) * entropy(subset) return
    total_entropy - weighted_entropy


# Find the best feature to split the data def
best_feature(data):
    features = data.columns[:-1] # Exclude the target column
    gains = {feature: information_gain(data, feature) for feature in features} return max(gains, key=gains.get)


# Create the decision tree
def id3(data, features=None):
    if len(data.iloc[:, -1].unique()) == 1: # All data points belong to the same class return
        data.iloc[:, -1].iloc[0]


    if len(features) == 0: # No more features to split on return
        data.iloc[:, -1].mode()[0]


    best = best_feature(data) tree
    = {best: {}}
```

```
    new_features = features.copy()
    new_features.remove(best)

    for value in data[best].unique(): subset
        = data[data[best] == value]
        tree[best][value] = id3(subset, new_features)

    return tree

# Function to classify new examples based on the decision tree def
classify(tree, example):
    if not isinstance(tree, dict):
    return tree feature =
    list(tree.keys())[0] value =
    example[feature]
    return classify(tree[feature][value], example)

# Function to visualize the decision tree using Graphviz
def create_tree_diagram(tree, dot=None, parent_name="Root", parent_value=""):
    if dot is None:
        dot = Digraph(format="png", engine="dot")
```

```
    if isinstance(tree, dict): # Tree node for
        feature, branches in tree.items():
            feature_name = f"{parent_name}_{feature}" dot.node(feature_name,
            feature)
            dot.edge(parent_name, feature_name, label=parent_value)


            for value, subtree in branches.items():
                value_name = f"{feature_name}_{value}"
                dot.node(value_name, f"{feature}: {value}")
                dot.edge(feature_name, value_name,
                label=str(value))
                # Recurse for each subtree create_tree_diagram(subtree,
    dot, value_name, str(value)) else: # Leaf node
        dot.node(parent_name + "_class", f"Class: {tree}") dot.ede(parent_name,
        parent_name + "_class", label="Leaf")
    return dot
# Example usage
data = pd.DataFrame({
    'Outlook': ['Sunny', 'Sunny', 'Overcast', 'Rain', 'Rain', 'Rain', 'Overcast', 'Sunny', 'Sunny', 'Rain',
'Sunny', 'Overcast', 'Overcast', 'Rain'],
    'Temperature': ['Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool', 'Cool', 'Mild', 'Cool', 'Mild', 'Mild', 'Mild', 'Hot',
'Mild'],
    'Humidity': ['High', 'High', 'High', 'High', 'High', 'Low', 'Low', 'High', 'Low', 'Low', 'Low', 'High', 'Low',
'High'],
```

```
    'Wind': ['Weak', 'Strong', 'Weak', 'Weak', 'Weak', 'Weak', 'Strong', 'Weak', 'Weak', 'Strong', 'Strong',
'Weak', 'Strong', 'Weak'],

    'PlayTennis': ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'No']

})

# Train the decision tree

tree = id3(data, features=list(data.columns[:-1])) print("Decision

Tree:", tree)

# Classify a new example

example = {'Outlook': 'Sunny', 'Temperature': 'Cool', 'Humidity': 'Low', 'Wind': 'Strong'} prediction

= classify(tree, example)

print("Prediction for the example:", prediction)



# Visualize the decision tree dot = create_tree_diagram(tree) dot.render("decision_tree",

view=True) # This will generate and open the tree diagram
```

Program 4

Implement Linear and Multi-Linear Regression algorithm for appropriate dataset

Screenshot:

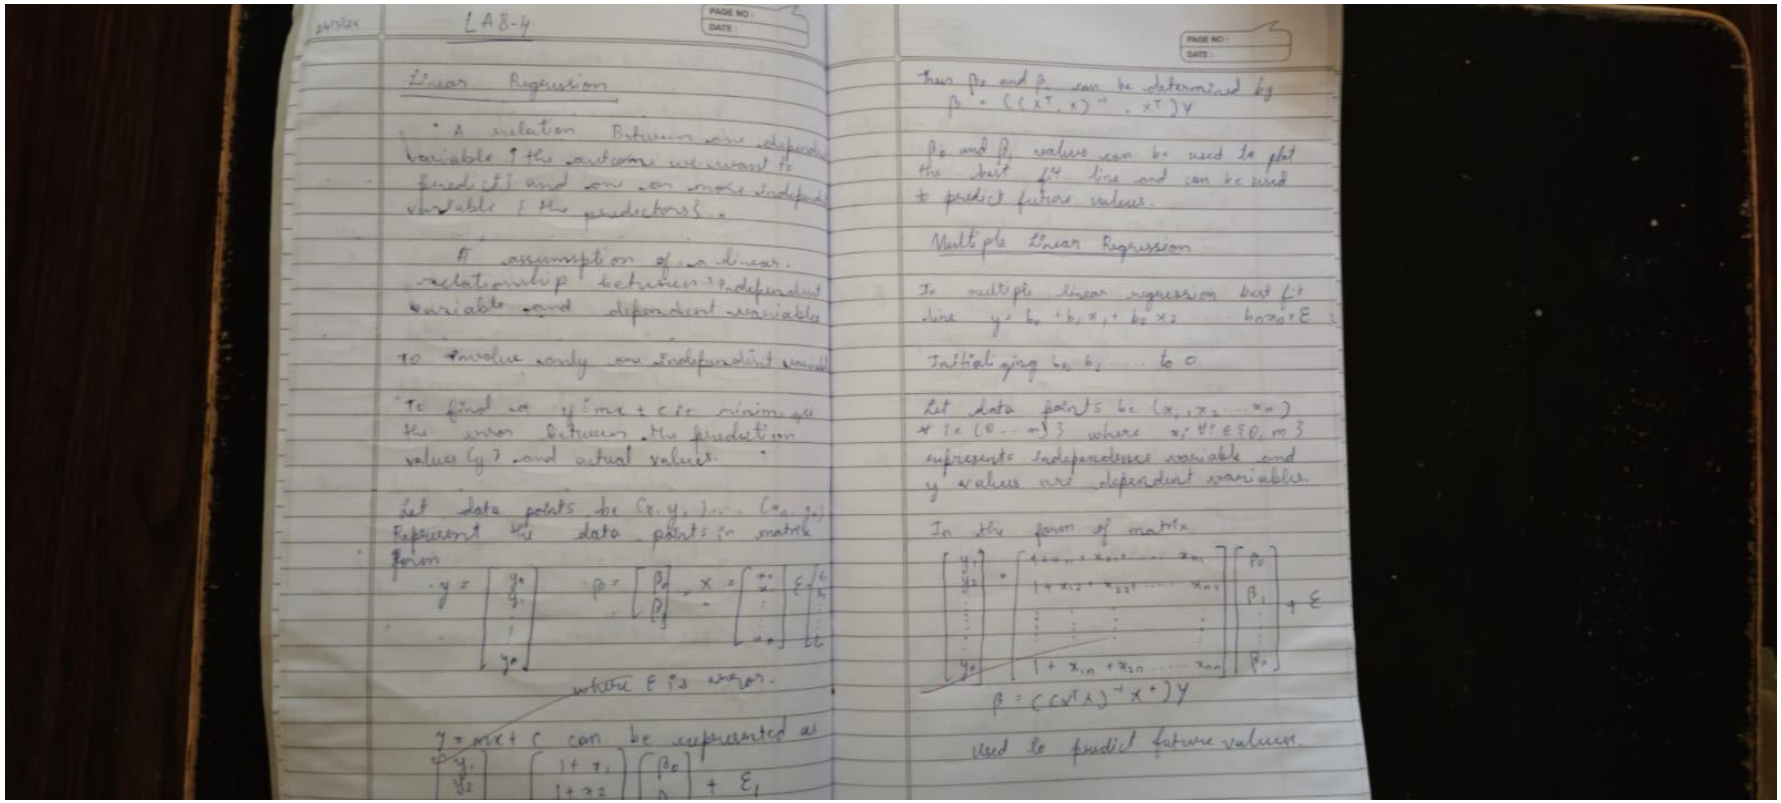

LAB-4

Linear Regression

• A relation between one dependent variable (the outcome we want to predict) and one or more independent variable (the predictors).

• assumption of a linear relationship between the dependent variable and independent variables

To involve only one independent variable

To find a $y = mx + c$ i.e. minimize the error between the prediction values (y) and actual values.

Let data points be $(x_1, y_1) \ldots (x_n, y_n)$

Represent the data points in matrix form

$$y = \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_n \end{bmatrix} \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} \quad X = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad \varepsilon \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix}$$

where ε is error.

$y = mx + c$ can be represented as

$$\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} \begin{bmatrix} 1 + x_1 \\ 1 + x_2 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \varepsilon_1$$

Then β_0 and β_1 can be determined by

$$\beta = ((X^T \cdot X)^{-1} \cdot X^T) Y$$

β_0 and β_1 values can be used to plot the best fit line and can be used to predict future values.

Multiple linear Regression

In multiple linear regression best fit line $y = b_0 + b_1 x_1 + b_2 x_2 \ldots b_n x_n + \varepsilon$

Initializing $b_0, b_1 \ldots b_n$ to 0

Let data points be $(x_1, x_2 \ldots x_n)$ where $x_i \forall i \in [0 \ldots m]$ where $x_i \forall i \in [0, m]$ represents independence variable and y values are dependent variable.

In the form of matrix

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 + x_{1,1} + x_{2,1} \ldots x_{n1} \\ 1 + x_{12} + x_{22} \ldots x_{n2} \\ \vdots \\ 1 + x_{1n} + x_{2n} \ldots x_{nn} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \end{bmatrix} + \varepsilon$$

$$\beta = (X^T X)^{-1} X^T) Y$$

Used to predict future values.

Code:

Linear Regression

```
import pandas as pd

df = pd.read_csv("/content/tvmarketing.csv") df


# Visualise the relationship between the features and the response using scatterplots

df.plot(x='TV',y='Sales',kind='scatter')


from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(df['TV'], df['Sales'], test_size=0.2, random_state=42)

from sklearn.linear_model import LinearRegression model = LinearRegression()
model.fit(x_train.values.reshape(-1, 1), y_train) y_train model.coef_
model.intercept_
```

MultiLinearRegression

```
import pandas as pd #

Step 2 : import data

house = pd.read_csv('https://github.com/YBIFoundation/Dataset/raw/main/Boston.csv')


# display first 5 rows
house.head()
```

```
y = house['MEDV']

X = house.drop(['MEDV'],axis=1)

# Step 4 : train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.7, random_state=2529)
# Step 5 : select model from sklearn.linear_model
import LinearRegression model =
LinearRegression() # Step 6 : train or fit model
model.fit(X_train,y_train) model.intercept_

model.coef_
```

Program 5

Build Logistic Regression Model for a given dataset

Screenshot:

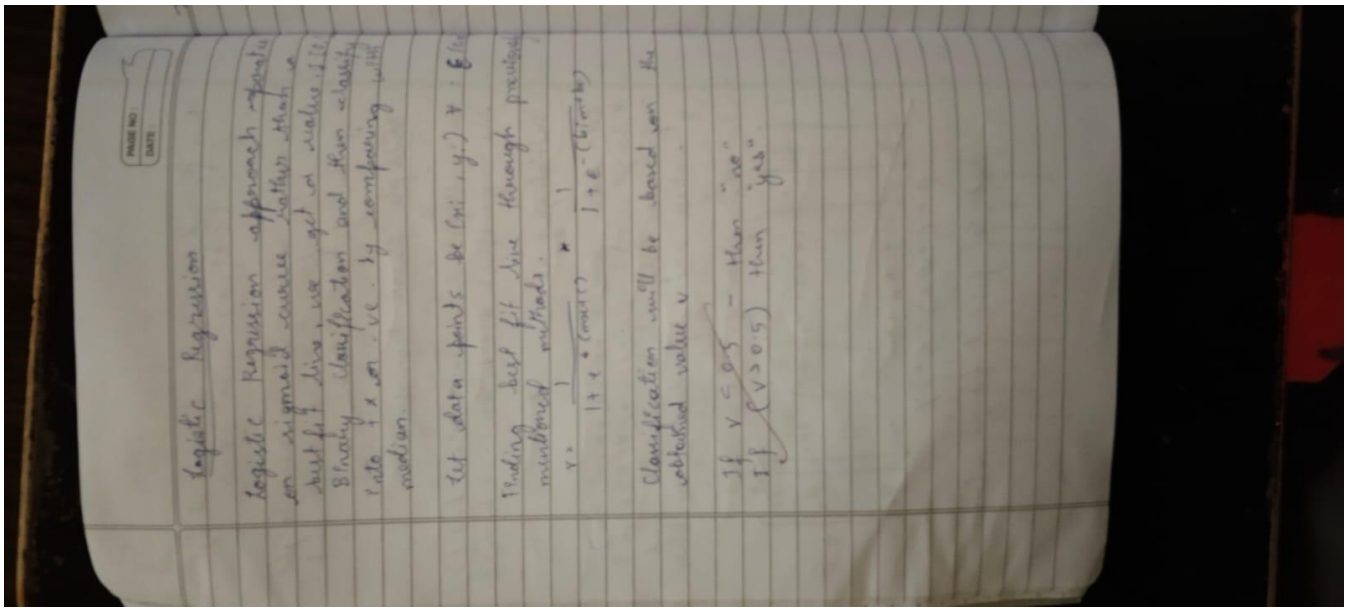

Code:

```
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris from
sklearn.model_selection import train_test_split from
sklearn.metrics import accuracy_score


# Load sample dataset (binary classification - Iris with only 2 classes)
iris = load_iris() X = iris.data[iris.target != 2] y = iris.target[iris.target
!= 2]


# Train/Test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
# Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Program 6

Build KNN Classification model for a given dataset

Screenshot:

7/4/25

Lab 5 KNN

PAGE NO :
DATE :

→ Input
Dataset $D = \{(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)\}$ Test data x.
Number of neighbours k

→ Output
Predicted labels. Used for test data
Classification accuracy.

Algorithm:

① Load the Dataset.
Input feature x
Target labels y

② Split the Dataset
Divide the dataset into:
Training set: X_{train}, y_{train}
Test set: X_{test}, Y_{test}

Use a fixed random seed for reproducibility

③ Initialized the KNN classifier
set the number of neighbours K.

④ Train the classifier
Store the trainer data X_{train}, y_{train}

⑤ predict for each Test Instance.

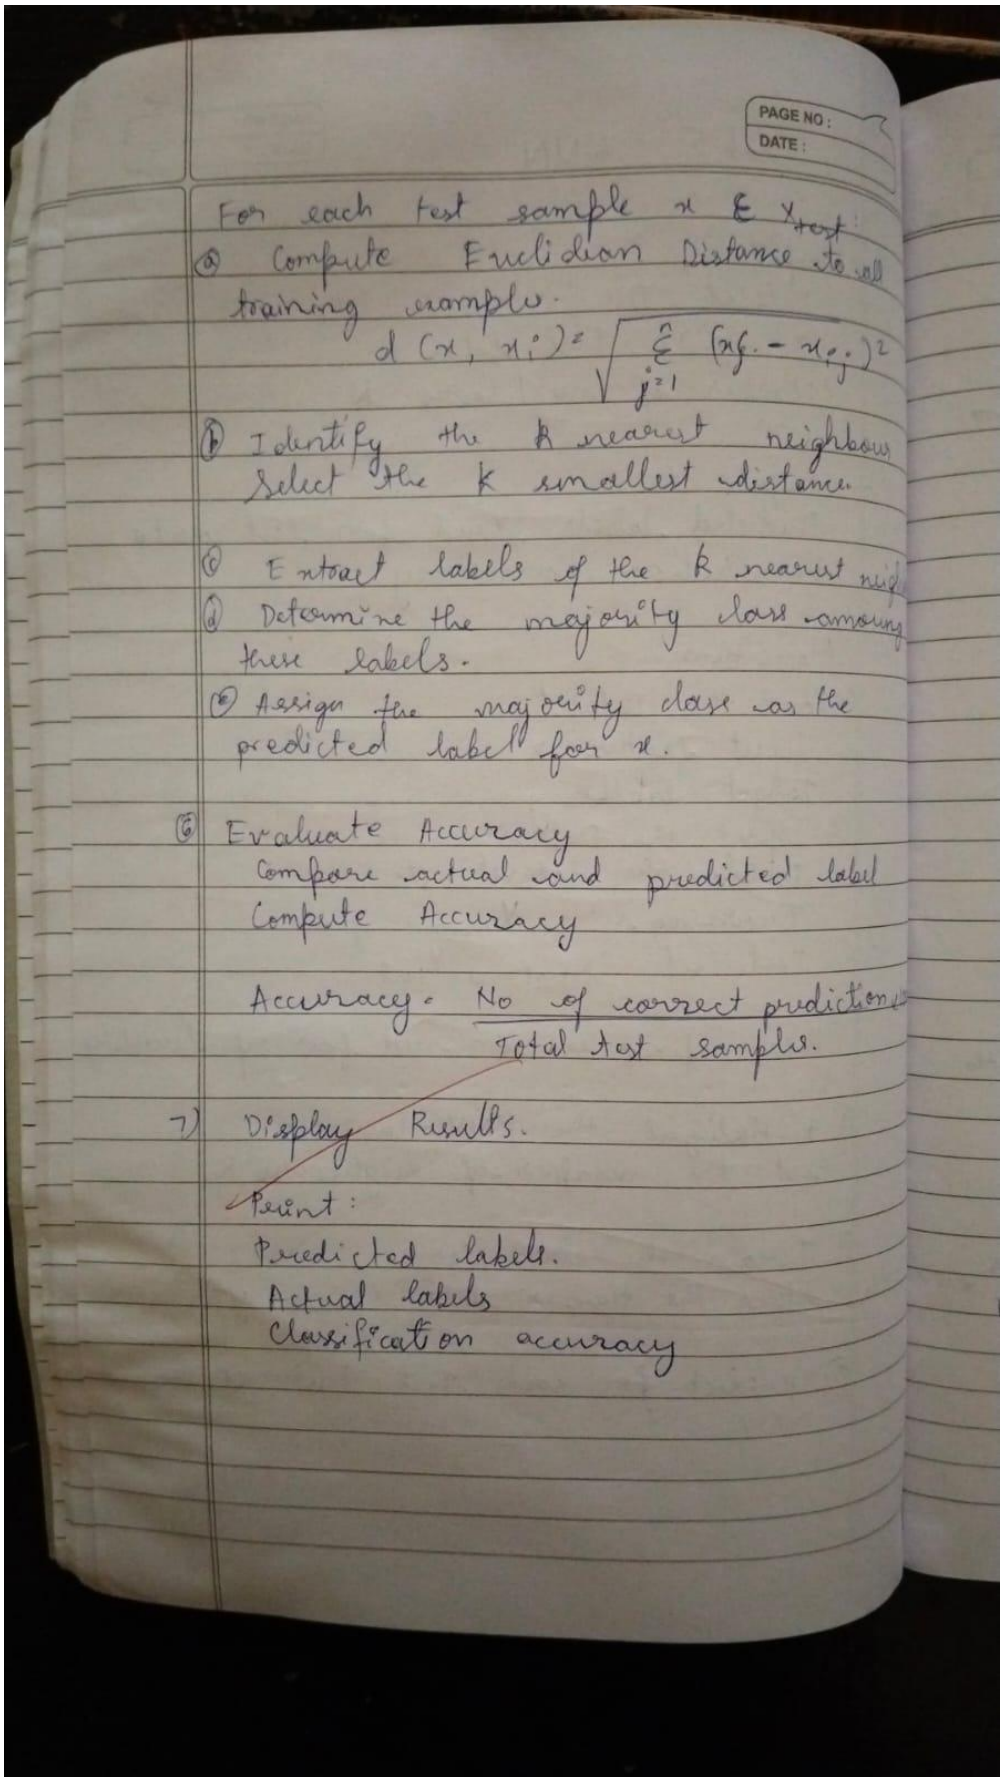

PAGE NO :
DATE :

For each test sample $x \in X_{test}$:

(a) Compute Euclidean Distance to all training samples.

$$d(x, x_i) = \sqrt{\sum_{j=1}^{n} (x_j - x_{ij})^2}$$

(b) Identify the k nearest neighbours. Select the k smallest distance.

(c) Extract labels of the k nearest neighbours.

(d) Determine the majority class among these labels.

(e) Assign the majority class as the predicted label for x.

(f) Evaluate Accuracy
Compare actual and predicted label
Compute Accuracy

Accuracy = $\frac{\text{No of correct predictions}}{\text{Total test samples.}}$

7) Display Results.

Print:
Predicted labels.
Actual labels
Classification accuracy

Code:

KNN

```
import numpy as np

from collections import Counter


class KNN:

def _init_(self, k=3): self.k = k

    def fit(self, X, y):

        self.X_train = np.array(X) self.y_train

        = np.array(y)


    def euclidean_distance(self, x1, x2):

        return np.sqrt(np.sum((x1 - x2) ** 2))


    def predict(self, X):

        predictions = [self._predict(x) for x in X] return

        np.array(predictions)


    def _predict(self, x):

        # Compute distances to all training points

        distances = [self.euclidean_distance(x, x_train) for x_train in self.X_train]
```

```
        # Get indices of k nearest neighbors
        k_indices = np.argsort(distances)[:self.k]


        # Get the labels of those neighbors

        k_nearest_labels = [self.y_train[i] for i in k_indices]

        # Return the most common label most_common =

        Counter(k_nearest_labels).most_common(1)          return

        most_common[0][0]


# Sample dataset (like a mini version of Iris) X_train

= [[1, 2], [2, 3], [3, 1], [6, 5], [7, 7], [8, 6]]

y_train = [0, 0, 0, 1, 1, 1]


# Test data

X_test = [[5, 5], [1, 1]]


# Using the KNN modelh knn

= KNN(k=3)

knn.fit(X_train, y_train)

predictions = knn.predict(X_test)


print("Predictions:", predictions)
```

Program 7

Build Support vector machine model for a given dataset

Screenshot:

PAGE NO :
DATE :

Support Vector Machine

- Input.

Dataset . $D = \{(x_1, y_1), (x_2, y_2) \ldots (x_n, y_n)\}$
Test dataset . S_{test} . Regularization parameters (Max no of iteration)

- Output

Predicted class labels ,
Classification accuracy

(1) Data loading and preprocessing
- Load the Iries dataset.
- Apply . Z-score normalization to standardize feature. $x' = \frac{x - \mu}{\sigma}$

(2) Split the Dataset.
- Split the data into
Training set (70) % , Test set (70 %)

(3) Initialize the SVM classifier
- Set the following parameter
C: Regularization constant.
$K(x, x') = x \cdot x'$

(4) One vs Rest Training Strategy
For each class C in the set of unique classes:
Convert labels into
Binary format.

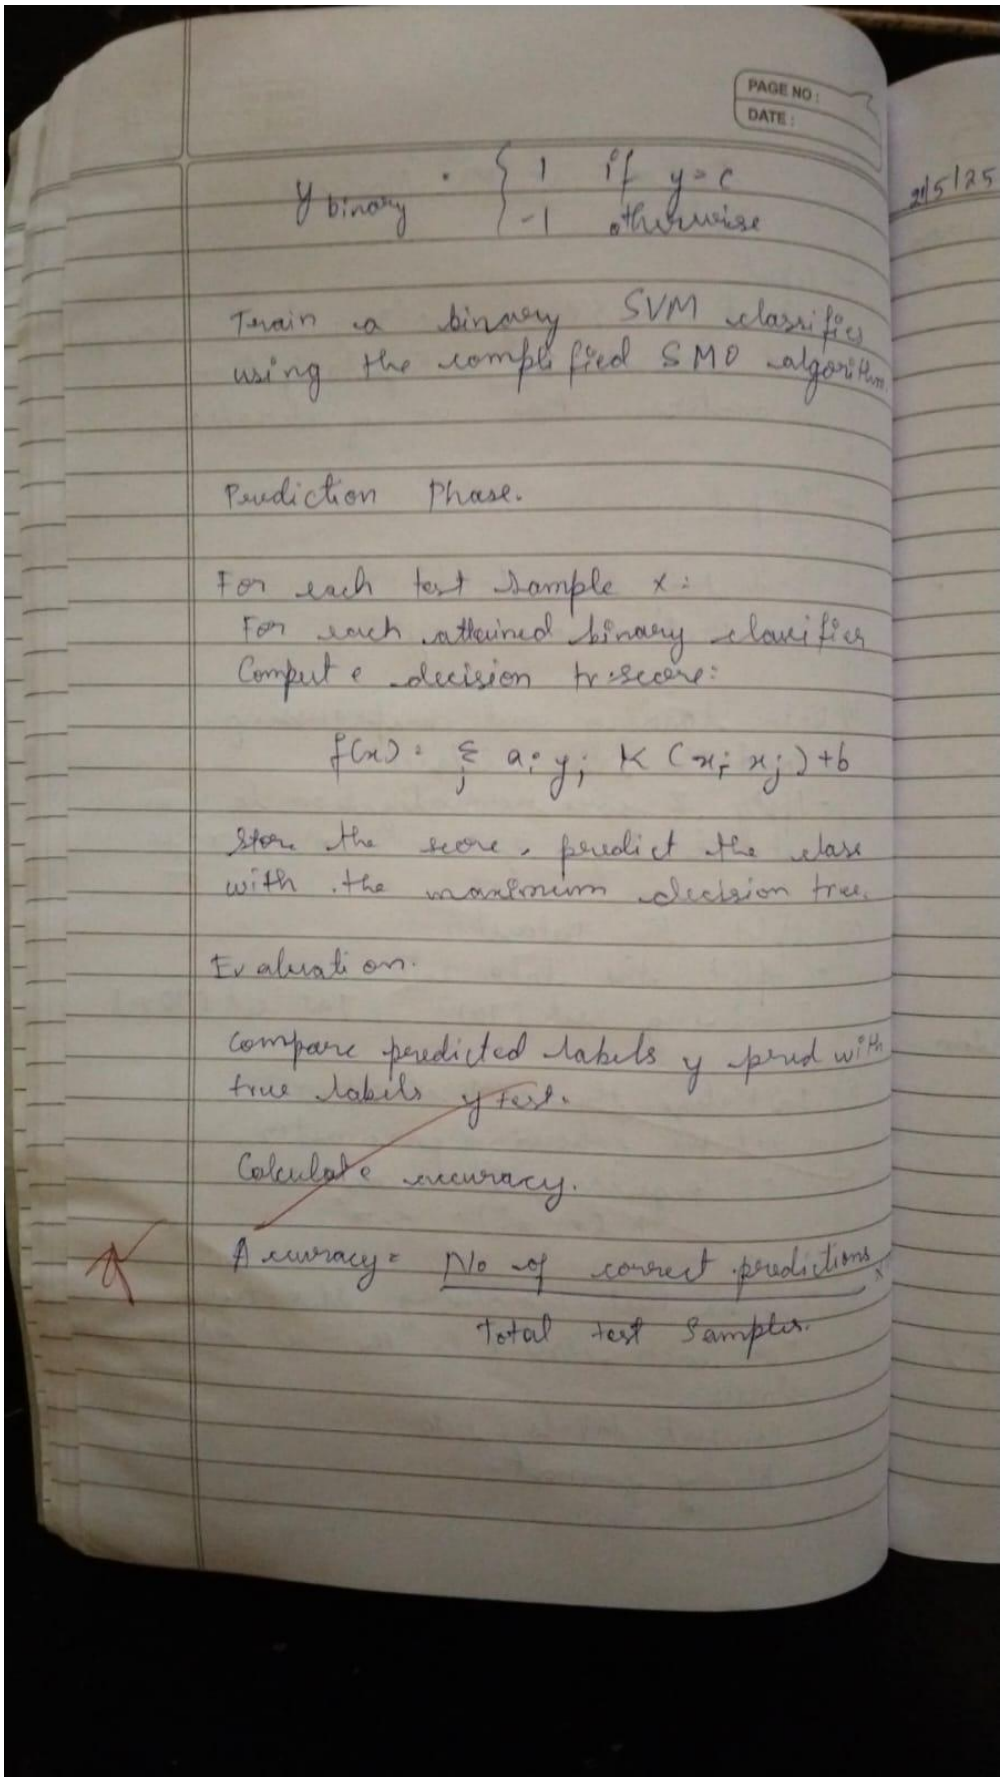

PAGE NO :
DATE :

$y_{binary} \cdot \begin{cases} 1 & \text{if } y = c \\ -1 & \text{otherwise} \end{cases}$

Train a binary SVM classifier using the simplified SMO algorithm.

Prediction Phase.

For each test sample x:
For each attained binary classifier
Compute decision function score:

$$f(x) = \sum_i a_i y_i K(x_i, x_j) + b$$

Store the score, predict the class with the maximum decision tree.

Evaluation.

Compare predicted labels y_{pred} with true labels y_{test}.

Calculate accuracy.

Accuracy = $\frac{\text{No of correct predictions}}{\text{Total test samples}}$

Code:

```
from sklearn import datasets
from sklearn.model_selection import train_test_split from sklearn.svm import SVC
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

# Load dataset
iris = datasets.load_iris()
X = iris.data y
= iris.target

# For binary classification (class 0 vs 1) X
= X[y != 2]
y = y[y != 2]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train SVM
clf = SVC(kernel='linear') # Try 'rbf', 'poly', etc.
clf.fit(X_train, y_train)
```

```
# Accuracy

print("Test Accuracy:", clf.score(X_test, y_test))
```

Program 8

Implement Random forest ensemble method on a given dataset

Screenshot:

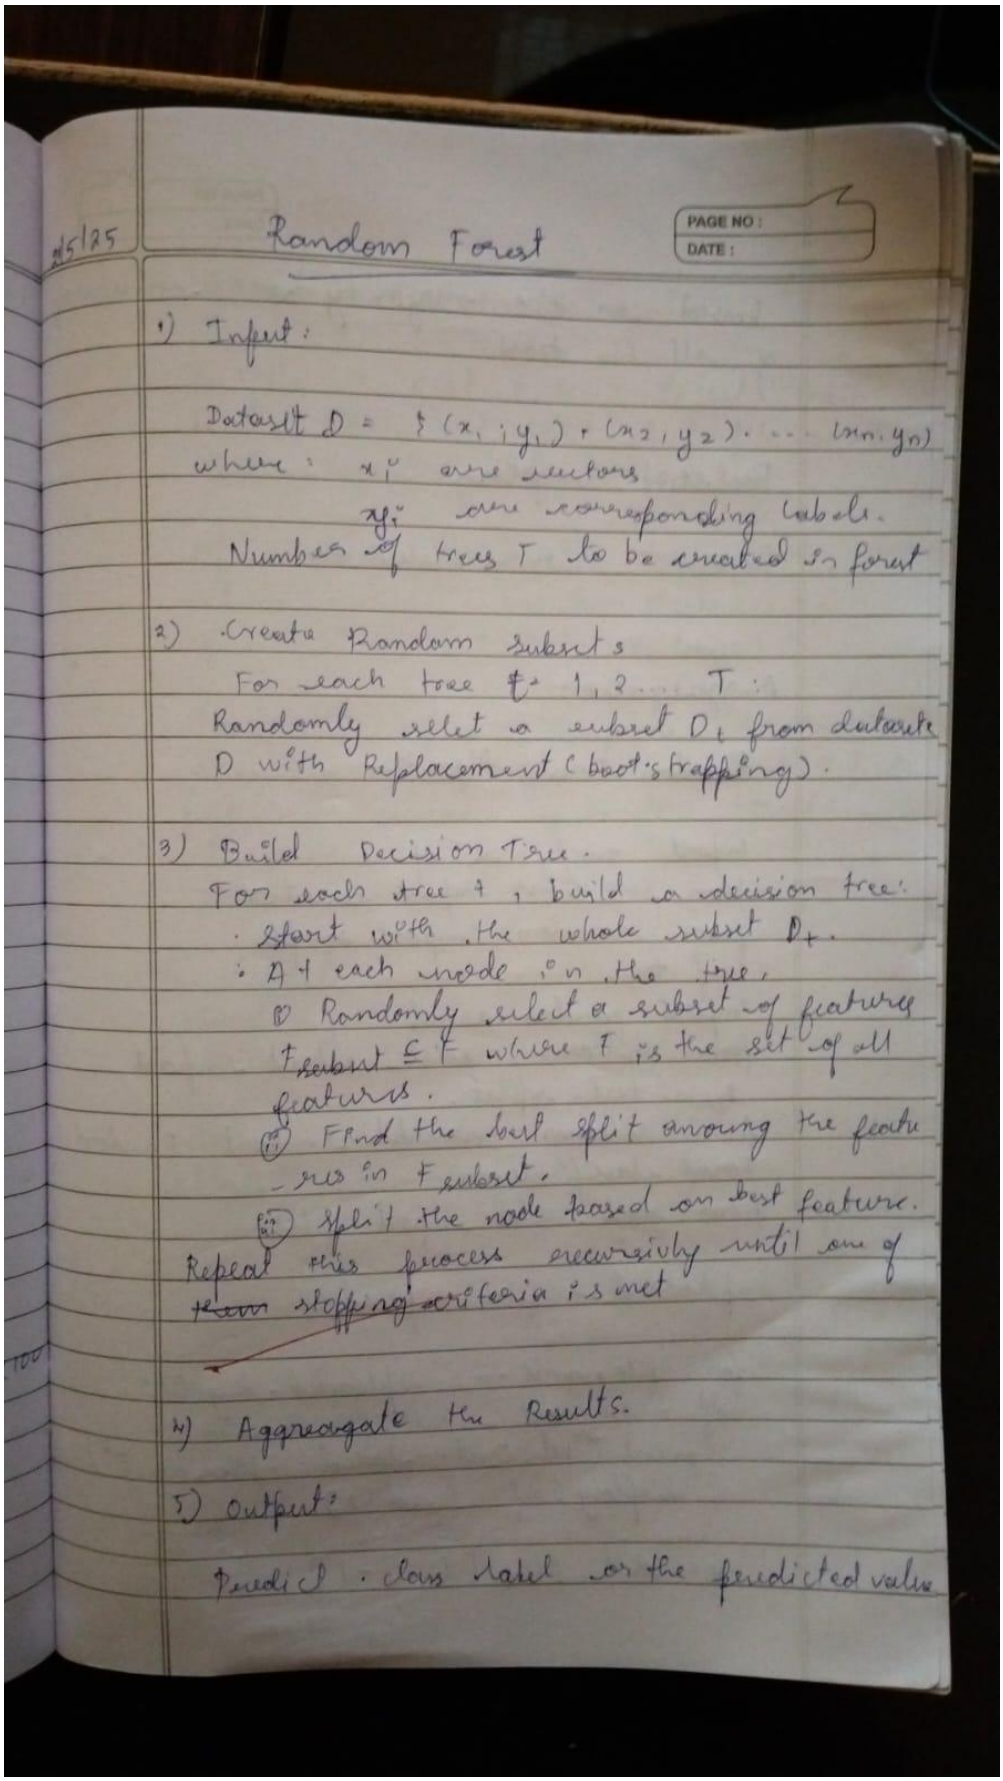

Random Forest

PAGE NO :
DATE :

1) Input:

Dataset $D = \{(x_1, y_1), (x_2, y_2), \ldots (x_n, y_n)\}$
where: x_i are vectors
y_i are corresponding labels.
Number of trees T to be created in forest

2) Create Random subsets
For each tree $t = 1, 2 \ldots T$:
Randomly select a subset D_t from dataset D with Replacement (bootstrapping).

3) Build Decision Tree.
For each tree t, build a decision tree:
- Start with the whole subset D_t.
- At each node in the tree,
  (i) Randomly select a subset of features $F_{subset} \subseteq F$ where F is the set of all features.
  (ii) Find the best split among the features in F_{subset}.
  (iii) Split the node based on best feature.

Repeat this process recursively until one of ~~the~~ stopping criteria is met

4) Aggregate the Results.

5) Output:

Predict: class label or the predicted value

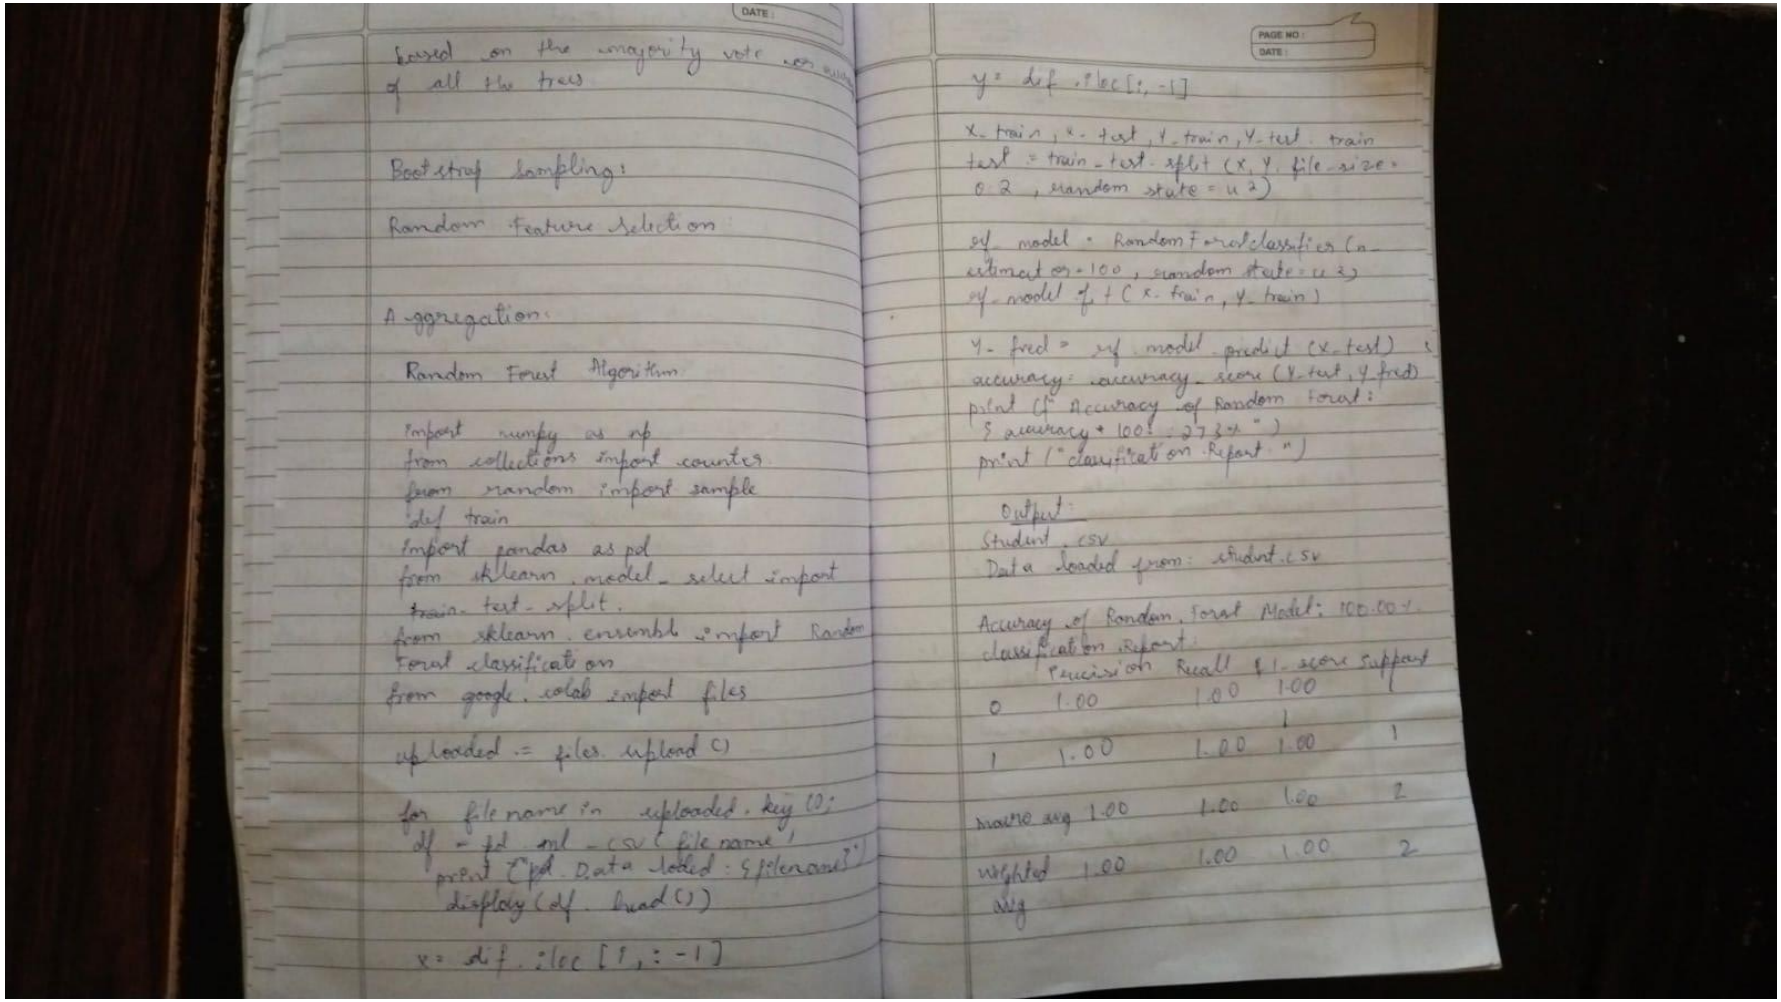

based on the majority vote on
of all the trees

Bootstrap Sampling:

Random Feature Selection:

Aggregation:

Random Forest Algorithm

```
import numpy as np
from collections import counter
from random import sample
def train
import pandas as pd
from sklearn.model_selection import
train_test_split
from sklearn.ensemble import Random
Forest classification
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():
  df = pd.read_csv(filename)
  print("pd.Data loded : {filename}")
  display(df.head())

x = df.iloc[:,:-1]
y = df.iloc[:,-1]

x_train, x_test, y_train, y_test = train
test = train_test_split(x, y, file_size=
0.2, random_state = 42)

rf_model = RandomForestclassifier(n_
estimators=100, random_state=42)
rf_model.fit(x_train, y_train)

y_pred = rf_model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy of Random Forest:
{accuracy*100:.2f}%")
print("classification Report:")
```

Output:

Student.csv
Data loaded from: student.csv

Accuracy of Random Forest Model: 100.00%
classification Report:

	Precision	Recall	f1-score	Support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2

Code:

```
from sklearn.datasets import load_iris
from  sklearn.model_selection  import  train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
# Load sample dataset iris
= load_iris()
X, y = iris.data, iris.target

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42) rf.fit(X_train,
y_train)

# Predict and evaluate y_pred = rf.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Program 9

Implement Boosting ensemble method on a given dataset

Screenshot:

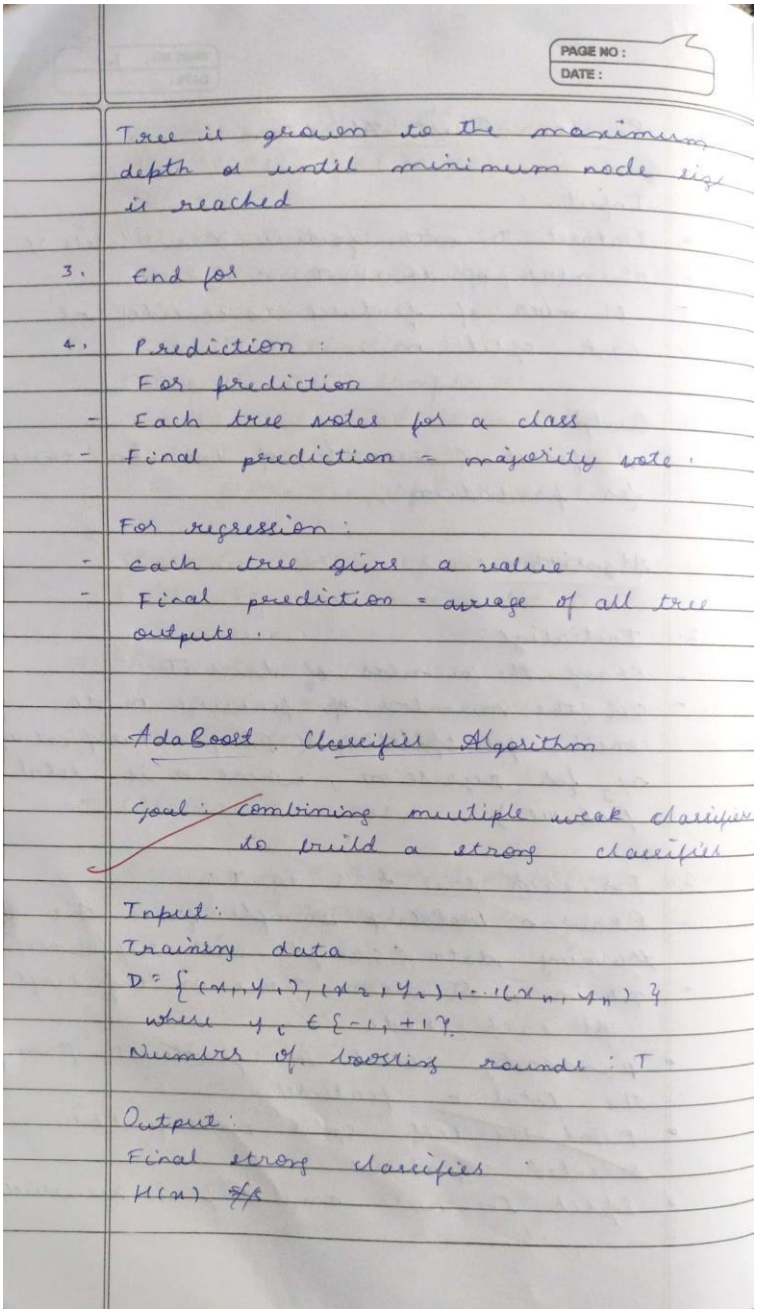

PAGE NO :
DATE :

Tree is grown to the maximum depth or until minimum node size is reached

3. End for

4. Prediction :
For prediction
- Each tree votes for a class
- Final prediction = majority vote.

For regression :
- Each tree gives a value
- Final prediction = average of all tree outputs.

AdaBoost Classifier Algorithm

Goal : Combining multiple weak classifiers to build a strong classifier

Input :
Training data
$D = \{(x_1, y_1), (x_2, y_2), \ldots (x_n, y_n)\}$
where $y_i \in \{-1, +1\}$
Number of boosting rounds : T

Output :
Final strong classifier :
$H(x)$

Code:

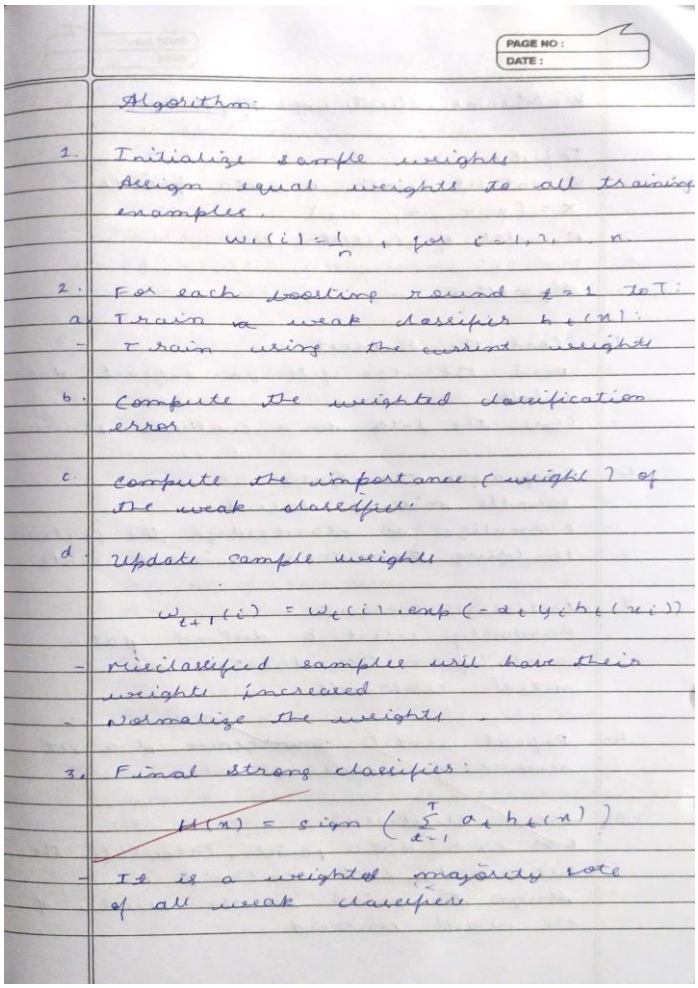

from sklearn.ensemble import AdaBoostClassifier from sklearn.datasets import load_iris from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score

Load Iris dataset iris

= load_iris()

X, y = iris.data, iris.target
For AdaBoost, we'll use binary classification

Convert to binary (setosa vs. not-setosa)

y = (y == 0).astype(int)

```
# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train AdaBoost
model = AdaBoostClassifier(n_estimators=50, learning_rate=1.0, random_state=42)
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
print("AdaBoost Accuracy (sklearn):", accuracy_score(y_test, y_pred))
```

Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file

Screenshot:

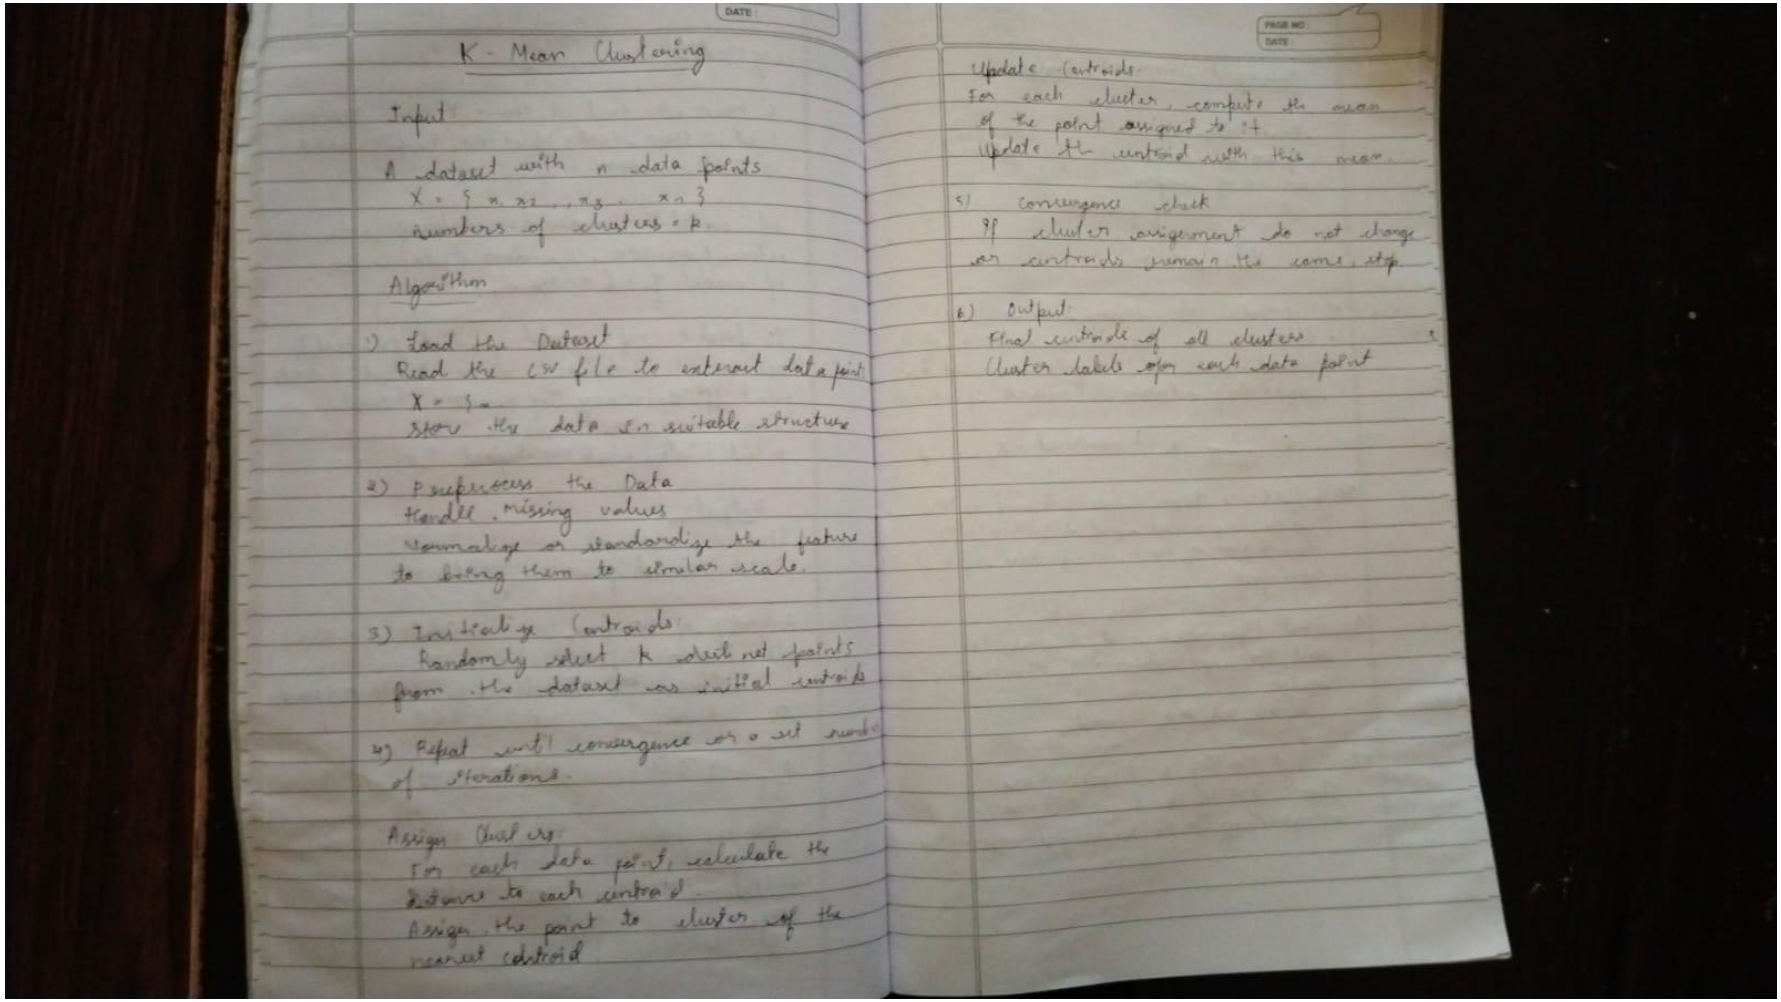

K - Mean Clustering

Input

A dataset with n data points
$X = \{x_1, x_2, \ldots, x_n\}$
numbers of clusters = k.

Algorithm

1) Load the Dataset
Read the CSV file to extract data point
$X = \{...\}$
Store the data in suitable structure

2) Preprocess the Data
Handle missing values
normalize or standardize the feature to bring them to similar scale.

3) Initialize Centroids
Randomly select k distinct points from the dataset as initial centroids

4) Repeat until convergence or a set number of iterations.

Assign Clusters:
For each data point, calculate the distance to each centroid.
Assign the point to cluster of the nearest centroid

Update Centroids:
For each cluster, compute the mean of the point assigned to it.
Update the centroid with this mean.

5) Convergence check
If cluster assignment do not change or centroids remain the same, stop

6) Output:
Final centroids of all clusters
Cluster labels for each data point

Code:

```
import pandas as pd

from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

from sklearn.datasets import load_iris # Import load_iris



# Step 1: Load the Iris dataset directly iris

= load_iris()

# Create a DataFrame from the data and target

data = pd.DataFrame(data=iris.data, columns=iris.feature_names)

# Add the target column for potential reference, though not used for clustering data['target']

= iris.target




# Step 2: Extract only numeric columns (or select required features)

# All features in the Iris dataset are numeric

X = data[iris.feature_names].values # Use the feature names to select columns



# Step 3: Apply KMeans

# Adjust n_clusters based on the expected number of clusters in your data (3 for Iris)

kmeans = KMeans(n_clusters=3, random_state=42, n_init=10) # Added n_init to suppress future warnings

data['Cluster'] = kmeans.fit_predict(X)
```

```
# Step 4: Plot clusters (for 2D data)
# Iris data has 4 features. We will plot the first two features for visualization. if
X.shape[1] >= 2:
    plt.scatter(X[:, 0], X[:, 1], c=data['Cluster'], cmap='viridis')
    plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], color='red', marker='x', s=200)
    plt.title("K-Means Clustering of Iris Dataset")
    plt.xlabel(iris.feature_names[0]) # Label with actual feature name
    plt.ylabel(iris.feature_names[1]) # Label with actual feature name
    plt.show()
else:
    print("Cannot plot clustering results directly for data with less than 2 features.")
```

Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method

Screenshot:

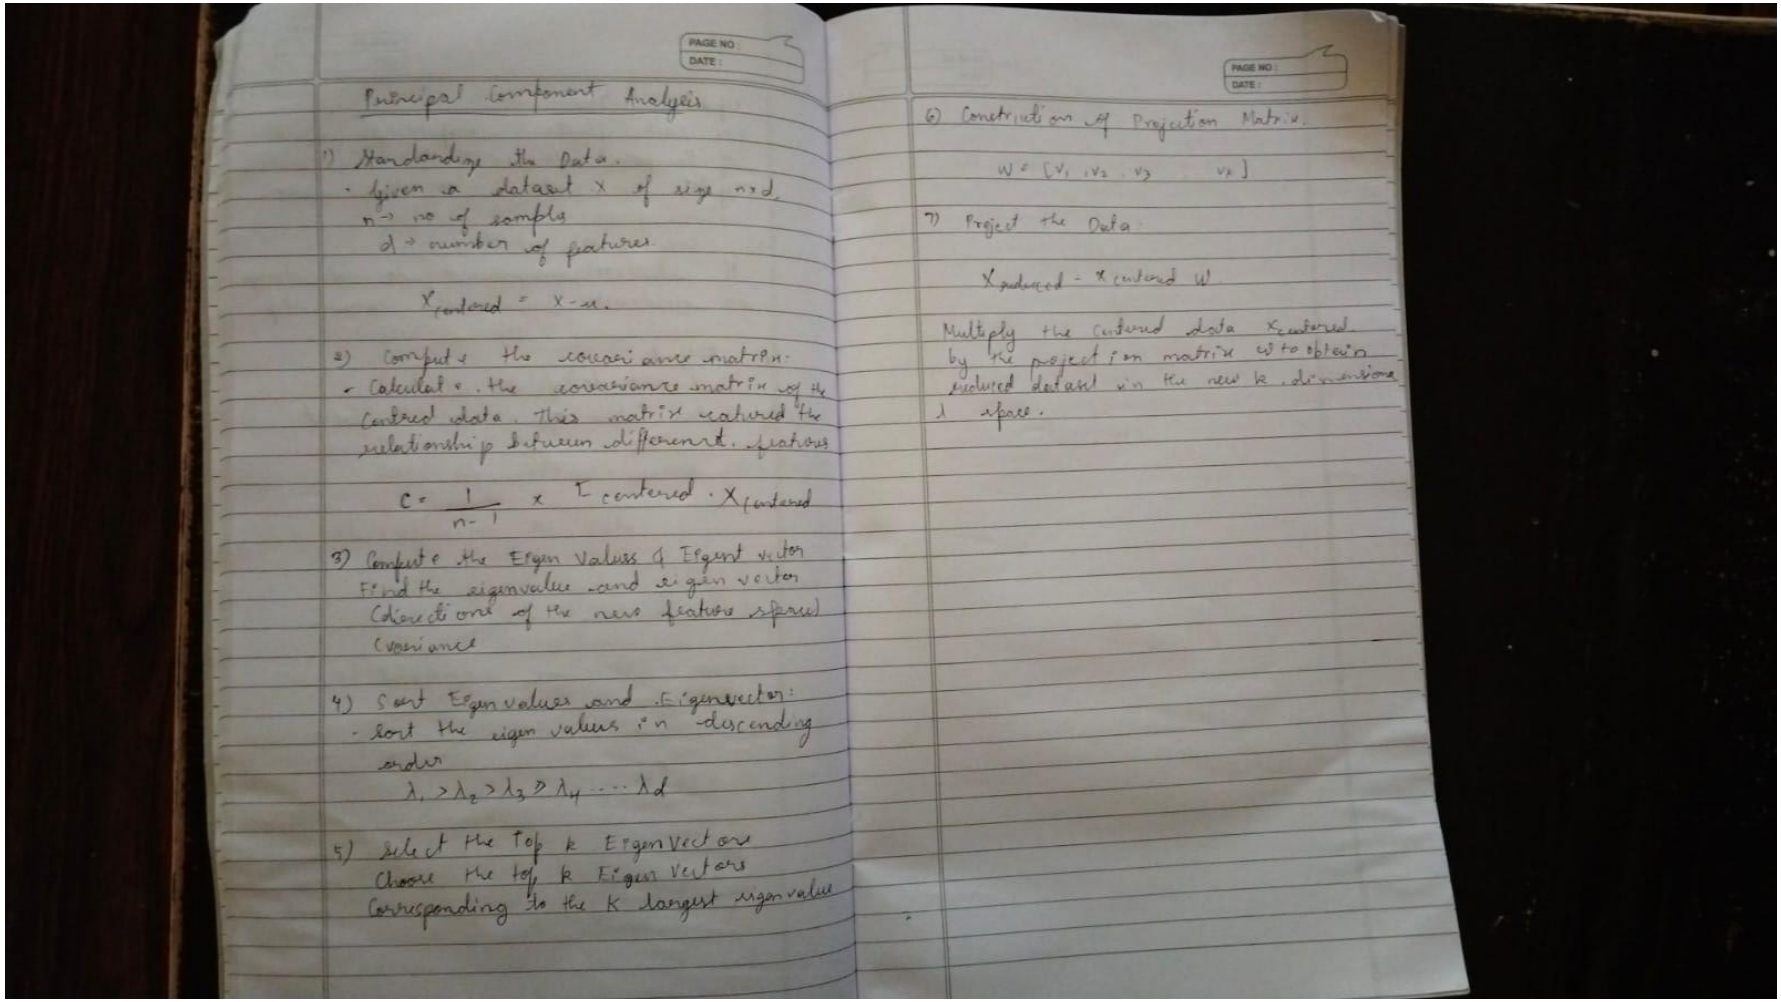

PAGE NO :
DATE :

Principal Component Analysis

1) Standardize the Data.
- Given a dataset X of size n×d,
n → no of samples
d → number of features.

$$X_{centered} = X - \mu.$$

2) Compute the covariance matrix.
- Calculate the covariance matrix of the centered data. This matrix captured the relationship between different features

$$C = \frac{1}{n-1} \times X_{centered}^{T} \cdot X_{centered}$$

3) Compute the Eigen Values & Eigen vector
Find the eigenvalue and eigen vector (directions) of the new feature space (covariance

4) Sort Eigen values and Eigenvector:
- Sort the eigen values in descending order

$$\lambda_1 > \lambda_2 > \lambda_3 > \lambda_4 \cdots \lambda_d$$

5) Select the Top k EigenVectors
Choose the top k Eigen Vectors corresponding to the K largest eigenvalue

PAGE NO :
DATE :

6) Construction of Projection Matrix.

$$W = [V_1, V_2, V_3 \quad V_k]$$

7) Project the Data

$$X_{reduced} = X_{centered} \cdot W$$

Multiply the centered data $X_{centered}$ by the projection matrix W to obtain reduced dataset in the new k dimensional space.

Code:

```
import pandas as pd

from sklearn.decomposition import PCA

from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt


# Load dataset

data = pd.read_csv("your_data.csv") # Replace with your file

X = data.select_dtypes(include=['float64', 'int64'])


# Step 1: Standardize scaler

= StandardScaler()

X_scaled = scaler.fit_transform(X)


# Step 2: Apply PCA

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X_scaled)


# Print explained variance ratio

print("Explained variance ratio:", pca.explained_variance_ratio_)


# Visualize
```

```
plt.scatter(X_pca[:, 0], X_pca[:, 1], c='blue', alpha=0.5) plt.title("PCA
- 2D Projection")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2") plt.show()
```

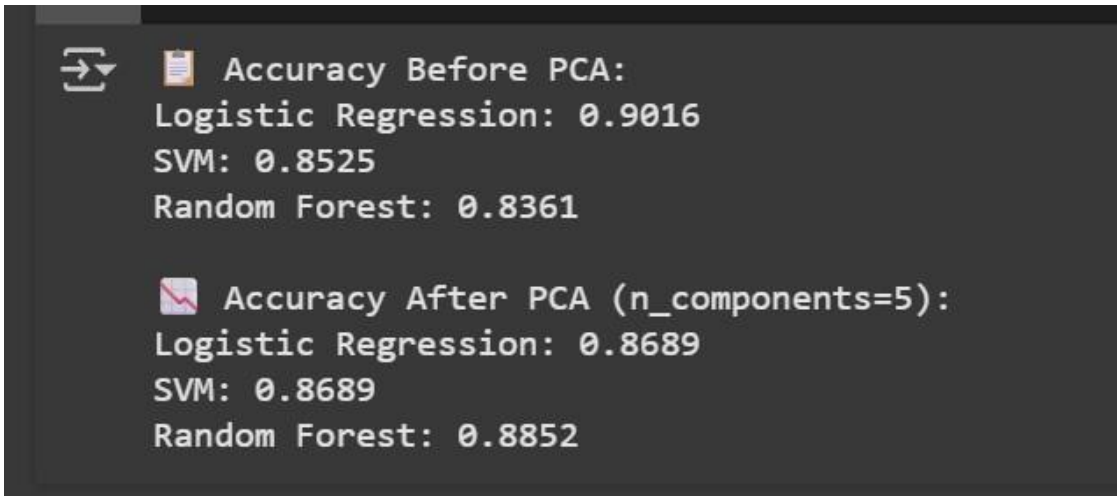
Accuracy Before PCA:
Logistic Regression: 0.9016
SVM: 0.8525
Random Forest: 0.8361
Accuracy After PCA (n_components=5):
Logistic Regression: 0.8689
SVM: 0.8689
Random Forest: 0.8852